



HyperionDev

Datasets and Dataframes

August 2024

Data Science Session Housekeeping

- The use of disrespectful language is prohibited in the questions, this is a supportive, learning environment for all - please engage accordingly.
- No question is daft or silly - **ask them!**
- There are **Q&A sessions** midway and at the end of the session, should you wish to ask any follow-up questions. Moderators are going to be answering questions as the session progresses as well.
- If you have any questions outside of this lecture, or that are not answered during this lecture, please do submit these for upcoming Academic Sessions. You can submit these questions here: [Questions](#)

Data Science Session Housekeeping cont.

- For all **non-academic questions**, please submit a query: www.hyperiondev.com/support
- Report a **safeguarding** incident: www.hyperiondev.com/safeguardreporting
- We would love your **feedback** on lectures: [Feedback on Lectures](#)

Learning Outcomes

- ❖ **Describe** the role and benefits of Python packages like Numpy and Pandas in data science tasks, including the use cases for Numpy arrays and Pandas dataframes.
- ❖ **Apply** Pandas functions to create and manipulate data frames from datasets, performing data exploration and transformation tasks.

Learning Outcomes

- ❖ **Create** and run a Jupyter Notebook to document and present data analysis workflows, including reading, exploring, and manipulating datasets using Pandas.

Problem Statement

Data science students are often required to handle, manipulate, and analyse large datasets efficiently. However, working with raw data using basic Python structures like lists and dictionaries can be cumbersome and inefficient, especially when dealing with complex operations or large volumes of data. Additionally, managing and presenting your work can be challenging without the right tools.

- ❖ How can we read datasets into our code more efficiently, and represent them using more appropriate data structures?
- ❖ How can we explore and manipulate our datasets using built-in functions in Python?

Lecture Overview

- Python Packages
 - ↳ Numpy
 - ↳ Pandas
- Dataframes in Pandas
- Jupyter Notebooks



Python packages for data science

NumPy



Python packages for data science: NumPy

❖ **NumPy** - stands for **Nu**merical **Py**thon

```
pip install numpy
```

❖ Numpy arrays differ from python lists

- Elements of same datatype
- Can handle arithmetic operations
- Preferred for larger chunks of data
- Requires proper modules to perform operations on them

❖ Array oriented programming, NumPy has smaller memory consumption, better runtime, and ease of data manipulation

Using NumPy

```
import numpy as np
```

```
import numpy as np

#Define a list
distances = [1, 13.1, 26.2, 100]
print(type(distances))
print(distances)
#Output: <class 'list'>
#Output: [1, 13.1, 26.2, 100]

#Convert to numpy array
numpy_dist = np.array(distances)
print(type(numpy_dist))
print(numpy_dist)
#Output: <class 'numpy.ndarray'>
#Output: [ 1.  13.1  26.2 100.]
```

Using NumPy

```
import numpy as np

#Define a List
distances = [1, 13.1, 26.2, 100]

#Convert to numpy array
numpy_dist = np.array(distances)

#Convert distances in miles to km
#Using numpy scalar mutliplication
conversion = numpy_dist * 1.60934
print(conversion)

#Output: [1.60934, 21.082354, 42.164708, 160.934]
```

```
#Using core Python

#Define a list
distances = [1, 13.1, 26.2, 100]

#Define an empty array to store km distances
conversion = []

#Using a for loop for conversion
for x in distances:
    conversion.append(x*1.60934)

print(conversion)

#Output: [1.60934, 21.082354, 42.164708, 160.934]
```

Python packages for data science

Pandas



Pandas

- ❖ **Pandas**, built on NumPy, is a Python module that contains high-level data structures and tools designed for fast and easy data analysis.

```
pip install pandas
```

```
import pandas as pd
```

- ❖ Fundamental pandas data structures

- Series (1-dimensional labelled array, can hold any data type)
- DataFrame (2-dimensional)
- Panel (pandas -> **panel data**)

	Name
0	Asha
1	Ben
2	Candice
3	Derek
4	Miriam
5	Seth
6	Zara

	Name	Age	Marks
0	Asha	12	96
1	Ben	12	92
2	Candice	13	94
3	Derek	12	96
4	Miriam	12	95
5	Seth	13	93
6	Zara	12	95

	DOB	Attending
0	1999-01-01	Yes
1	1999-01-01	Yes
2	1999-01-01	Yes
3	1999-01-01	Yes
4	1999-01-01	Yes
5	1999-01-01	Yes
6	1999-01-01	Yes

	Gender	Dietary	Location
0	Female	Vegetarian	London
1	Male	Vegetarian	London
2	Female	Vegetarian	London
3	Male	Vegetarian	London
4	Female	Vegetarian	London
5	Male	Vegetarian	London
6	Female	Vegetarian	London

	Name	Age	Marks
0	Asha	12	96
1	Ben	12	92
2	Candice	13	94
3	Derek	12	96
4	Miriam	12	95
5	Seth	13	93
6	Zara	12	95

Datasets



Datasets

- ❖ A dataset is a **structured collection of information** relevant to a specific investigation or project
- ❖ In data science, they provide the **raw material** for **analysis** and **modeling**. Understanding different dataset formats ensures you can work with data from various sources (databases, online repositories, etc.).
- ❖ With the help of **Pandas DataFrames**, we can effortlessly manipulate data to suit our needs.

DataFrames



DataFrames

- ❖ A DataFrame is the way the **Pandas library** in Python represents **tabular data**. It's like a powerful *spreadsheet* within your code.
- ❖ **Rows**: Each row represents a **single observation** or data point (e.g., a person, a product, a transaction).
- ❖ **Columns**: Each column represents a **variable or feature** (e.g., height, price, date). Data within a column usually shares the **same data type** (numbers, text, etc.).

Pandas DataFrame



Pandas DataFrame

- ❖ The pandas' library documentation defines a DataFrame as a “two-dimensional, size-mutable, with labelled rows and columns.”

```
import pandas
```

columns axis=1		column name	more columns to display							
		color	director_name	num_critic_for_reviews	duration	...	actor_2_facebook_likes	imdb_score	aspect_ratio	movie_facebook_likes
index label	0	Color	James Cameron	723.0	178.0	...	936.0	7.9	1.78	33000
	1	Color	Gore Verbinski	302.0	169.0	...	5000.0	7.1	2.35	0
	2	Color	Sam Mendes	602.0	148.0	...	393.0	6.8	2.35	85000
	3	Color	Christopher Nolan	813.0	164.0	...	23000.0	8.5	2.35	164000
	4	NaN	Doug Walker	NaN	NaN	...	12.0	7.1	NaN	0

index
axis=0

missing values

data
(values)

Anatomy of a DataFrame

Pandas DataFrame

- ❖ Pandas provides functions like `pd.read_csv()`, `pd.read_excel()`, `pd.read_sql()`, to bring your data directly into your coding environment as DataFrames.
- ❖ This is where you start turning your raw data into something easily workable.

```
import pandas as pd
```

```
# url = 'https://raw.githubusercontent.com/mwaskom/seaborn-data/master/iris.csv'  
# df = pd.read_csv(url)
```

```
iris = datasets.load_iris()
```

```
df = pd.DataFrame(iris.data, columns=iris.feature_names)
```


Exploring datasets

- ❖ **df.head(), df.tail():** Peek at the top and bottom rows for initial understanding

```
df.head()
```

✓ 0.0s

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

Exploring datasets

- ❖ **df.head(), df.tail()**: Peek at the top and bottom rows for initial understanding

```
df.tail()
```

✓ 0.0s

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species
145	6.7	3.0	5.2	2.3	2
146	6.3	2.5	5.0	1.9	2
147	6.5	3.0	5.2	2.0	2
148	6.2	3.4	5.4	2.3	2
149	5.9	3.0	5.1	1.8	2

Exploring datasets

- ❖ **df.shape:** Tells you the dimensions (rows, columns) of your data.

```
df.shape
```

```
✓ 0.0s
```

```
(150, 5)
```

Exploring datasets

- ❖ **df.info():** Gives the **data types** of each column, and if columns have missing values

```
df.info()
✓ 0.0s
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	sepal length (cm)	150 non-null	float64
1	sepal width (cm)	150 non-null	float64
2	petal length (cm)	150 non-null	float64
3	petal width (cm)	150 non-null	float64
4	species	150 non-null	int64

dtypes: float64(4), int64(1)
memory usage: 6.0 KB

Exploring datasets

- ❖ **df.describe():** Quick summary statistics for numerical columns.

```
df.describe()
```

✓ 0.0s

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	species
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333	1.000000
std	0.828066	0.435866	1.765298	0.762238	0.819232
min	4.300000	2.000000	1.000000	0.100000	0.000000
25%	5.100000	2.800000	1.600000	0.300000	0.000000
50%	5.800000	3.000000	4.350000	1.300000	1.000000
75%	6.400000	3.300000	5.100000	1.800000	2.000000
max	7.900000	4.400000	6.900000	2.500000	2.000000

Manipulating Data



HyperionDev

Manipulating Data

- ❖ **Selecting Columns:** You often work with a **subset of features**.
- ❖ Using `df[['column1', 'column2']]` gets you only specific columns.

```
df.columns
```

```
✓ 0.0s
```

```
Index(['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)',  
      'petal width (cm)', 'species'],  
      dtype='object')
```

```
# Select specific columns
```

```
df_selected = df[['species', 'petal length (cm)', 'petal width (cm)']]
```

```
✓ 0.0s
```

Manipulating Data

- ❖ **Filtering Rows:** Focus on specific subsets meeting certain conditions, e.g., `df[df['species'] == 'setosa']`

```
# Filter by flower species
```

```
df_setosa = df[df['species'] == 'setosa']
```

✓ 0.0s

Manipulating Data

- ❖ **Creating New Columns:** Derived features, e.g., calculating area from length and width.

```
# Create a new calculated column
```

```
df['petal area (cm^2)'] = df['petal length (cm)'] * df['petal width (cm)']
```

```
✓ 0.0s
```

Manipulating Data

- ❖ **Renaming/Dropping:** Improve clarity or get rid of unneeded data.

```
# Rename a column
```

```
df = df.rename(columns={'sepal length (cm)': 'sepal_len'})
```

```
✓ 0.0s
```

- ❖ Data manipulation gives you a **highly customized DataFrame** focused on your exact analysis needs.

Built-in Methods

- ❖ Pandas offers a toolbox of functions for calculations:
 - **mean()** - Computes the mean for each column.
 - **min()** - Computes the minimum for each column.
 - **max()** - Computes the maximum for each column.
 - **std()** - Computes the standard deviation for each column.
 - **var()** - Computes the variance for each column.
 - **unique()** - Computes the number of unique values in each column.
- ❖ This is the start of understanding the characteristics of your data.

Grouping and Aggregation

- ❖ `df.groupby()`: Divide your data **based on categories** in a column (e.g., group by species).

```
print(df['petal area (cm^2)'].mean())  
print(df['species'].nunique())  
print(df.groupby('species')['petal length (cm)'].std())
```

✓ 0.0s

5.794066666666667

3

species

0 0.173664

1 0.469911

2 0.551895

Name: petal length (cm), dtype: float64

Grouping and Aggregation

- ❖ **.agg()**: Apply calculations within each group (e.g., average length, maximum width).

```
df.groupby('species').agg(  
    mean_petal_length=('petal length (cm)', 'mean'),  
    max_sepal_width=('sepal width (cm)', 'max')  
)
```

✓ 0.0s

	mean_petal_length	max_sepal_width
species		
0	1.462	4.4
1	4.260	3.4
2	5.552	3.8

Jupyter Notebook



Jupyter Notebook

- ❖ An **interactive environment** perfect for **data science work**. They let you combine **code**, the **results of the code** (output), and **explanatory text** (like in a scientific report).
- ❖ This fosters **clear data exploration** and **storytelling**, all in one place

Installation

```
pip install jupyter
```

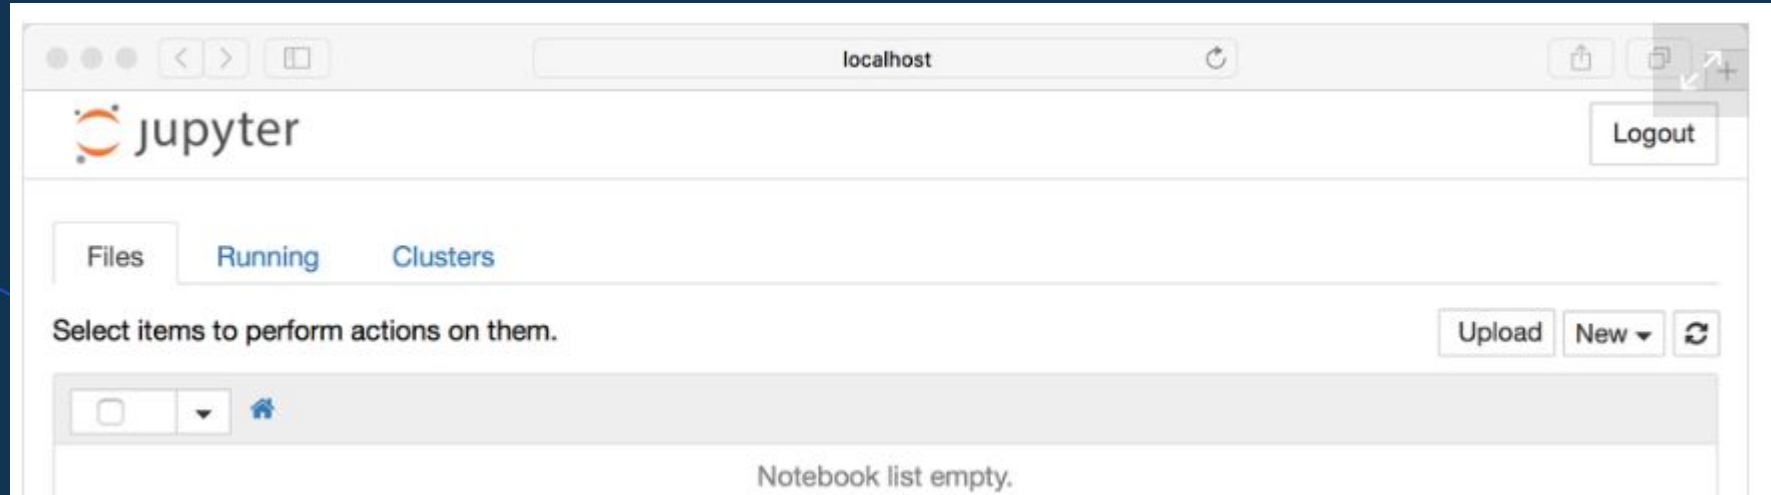
Running

```
jupyter notebook
```

```
python -m notebook
```

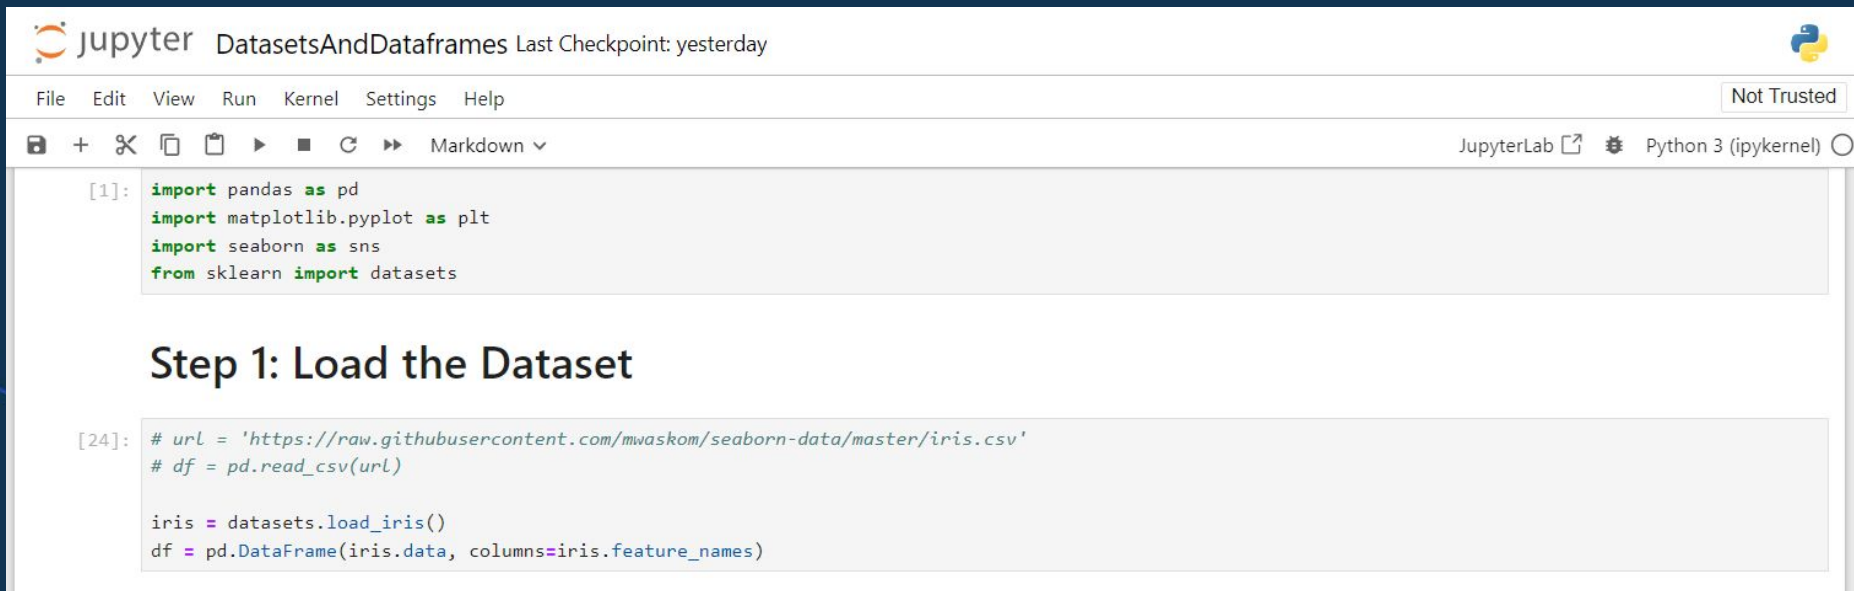
Jupyter Notebook

Starting the Jupyter Notebook Server, default browser goes to the URL <http://localhost:8888/tree>



Jupyter Notebook

Creating, naming notebook, code, markdown



The screenshot shows a Jupyter Notebook window titled "DatasetsAndDataframes" with a "Last Checkpoint: yesterday" status. The interface includes a menu bar (File, Edit, View, Run, Kernel, Settings, Help) and a toolbar with icons for saving, adding, deleting, and running code. A "Not Trusted" warning is visible in the top right. The notebook contains two cells: a code cell with imports for pandas, matplotlib, seaborn, and sklearn datasets, and a markdown cell titled "Step 1: Load the Dataset" containing code to load the Iris dataset from a GitHub URL.

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import datasets
```

Step 1: Load the Dataset

```
[24]: # url = 'https://raw.githubusercontent.com/mwaskom/seaborn-data/master/iris.csv'
# df = pd.read_csv(url)

iris = datasets.load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
```


Questions and Answers



Thank you for attending

